

Our Progress

- The progress so far
 - We saw some data structures: arrays, stacks, queues, linked lists, and hash tables.
 - Saw applications of the above data structures too.
 - Understand parallelism in computing to a small extent.
- This topic:
 - Study a framework for writing simple parallel programs.
 - Write simple parallel programs.

The OpenMP Programming Framework

- OpenMP is a specification for a set of compiler directives, library routines, and environment variables that can be used to specify shared memory parallelism in Fortran and C/C++ programs.
- The MP in OpenMP stands for Multi Processing.
- OpenMP is designed for Fortran, C and C++ to support the language that the underlying compiler supports.

A Quick Example

```
#include <stdio.h>
#include <omp.h>
/* Main Program */
main() {
    int ThreadID, NoofThreads;
    Omp_set_num_threads(3);
    /* OpenMP Parallel Directive */
    #pragma omp parallel private(ThreadID)
    {
        ThreadID = omp_get_thread_num();
        printf("Hello World is being printed by the thread id %d\n", ThreadID);
        /* Master Thread Has Its ThreadID 0 */
        if (ThreadID == 0) {
            printf("Master prints Numof threads");
            NoofThreads = omp_get_num_threads();
            printf("Total number of threads are %d", NoofThreads);
        }
    }
}
```

OpenMP Basics

- How to choose the number of threads to execute?
- Typically, depends on the machine used.
 - Plus other factors such as the nature of the program.
- How to set the number of threads?
- Several ways.
 1. Via environment variables.
 - `setenv OMP_NUM_THREADS 3` (if using `csh/tcsh`)
 - `export OMP_NUM_THREADS=3` (if using `bash/ksh`)

OpenMP Basics

- Another way is via the program.
 - **omp_set_num_threads**(*int num_threads*)
 - Allows one to change this number in between the program.

How to Compile and Run?

- The compilation and execution details of an OpenMP program may vary on different architectures.
- To compile a C-based OpenMP program:
 - **# gcc -o <Name of executable> <program name> -fopenmp**
 - In general, use:
 - # <compiler> -o <executable-name> <program-name>
 <compiler-flag> -lm
- To execute:
 - **./< Name of executable>**

Example Program Output

Hello World from thread = 0

Number of threads = 3

Hello World from thread = 2

Hello World from thread = 1

- The actual output you obtain may vary but will have the above model.

More Serious Examples

- Recall our prefix sum parallel algorithm reproduced below.

Input: Array A of size $n = 2k$

Output: Prefix sums in $C[0,j]$, $1 < j < n$

begin

1. for $1 < j < n$ *pardo*

$B[0,j] := A[j]$

2. for $h = 1$ to $\log n$ do

3. for $1 < j < n/2^h$ *pardo*

$B[h,j] := B[h-1,2j-1] + B[h-1,2j]$

4. for $h = \log n$ to 0 do

5. for $1 < j < n/2^h$ *pardo*

if j is even then $C[h,j] := C[h+1,j/2]$

else if $j = 1$ then $C[h,1] := B[h,1]$

else $C[h,j] := C[h+1,(j-1)/2] + B[h,j]$

end

How to Write a Parallel Program?

- OpenMP lets us write for-loops to execute in parallel.
- This is done by using **#pragma omp for**.
- A simple example first.

A Simpler Example

```
#pragma omp parallel \  
    shared(n,x,y) private(i)  
{  
    #pragma omp parallel for  
    for (i=0; i<n; i++)  
        x[i] += y[i];  
} /*-- End of parallel region --*/
```

- This program adds the *i*th element of *Y* to the *i*th element of *X*.

A Simpler Example

```
#pragma omp parallel \  
    shared(n,x,y) private(i)  
{  
    #pragma omp parallel for  
    for (i=0; i<n; i++)  
        x[i] += y[i];  
} /*-- End of parallel region --*/
```

- This program adds the *i*th element of *Y* to the *i*th element of *X*.
- What are these *shared* and *private*?
 - Indicate which data items are shared across threads, and which are not shared.
 - In this case, the arrays are to be shared by the threads, but the loop index should not be.

Model Execution

- Suppose we have $n = 10$ and we have 2 threads.
- Let $X = (1, 2, 3, \dots, 10)$ and $Y = (2, 3, 5, 7, 11, 13, 17, 19, 21, 23)$.

Thread 1 executes for
indices $i = 1, 2, 3, 4, 5$.

$i = 1: x[1] = 3$

$i = 2: x[2] = 5$

$i = 3: x[3] = 8$

$i = 4: x[4] = 11$

$i = 5: x[5] = 16$

Thread 2 executes for
indices $i = 6, 7, 8, 9, 10$.

$i = 6: x[6] = 19$

$i = 7: x[7] = 24$

$i = 8: x[8] = 27$

$i = 9: x[9] = 30$

$i = 10: x[10] = 33$

Back to Our Program

Begin

#pragma omp parallel for

shared(A, B, n), private(j)

1. for $1 < j < n$ *pardo*

$B[0,j] := A[j]$

Back to Our Program

Begin

```
#pragma omp parallel for  
    shared(n, A, B), private(j)  
1. for 1 < j < n pardo  
     $B[0,j] := A[j]$ 
```

- At the end of this statement, we wish to write the next part of the program.
- This corresponds to the following

Back to Our Program

```

Begin
Omp_set_num_threads(4)
#pragma omp parallel for
    shared(A, B, n), private(j)
    1. for (j=1; j<=n; j++)
        B[0,j] := A[j]
    2. . for (h = 1; h<=log n; h++)
        Omp_set_num_threads(8)
        #pragma omp parallel for
            shared(B, n), private(j)
            3. for 1 < j < n/2h pardo
                B[j] = B[h-1, 2j-1] + B[h-1, 2j]

```

- At the end of this statement, we wish to write the next part of the program.
- This corresponds to the following

Back to Our Example

- The pragma is not on the outer loop.
- The outer loop has to be done in sequence.
- The inner loop can be done in parallel.

Back to Our Example

- The third loop has a similar property. The outer loop has to run in sequence.

```
for  $h = \log n$  to 0 do  
  #pragma omp parallel for shared(C, B) private(j,n,h)  
  5. for  $1 < j < n/2^h$   
    if  $j$  is even then  $C[h,j] := C[h+1,j/2]$   
    else if  $i = 1$  then  $C[h,1] := B[h,1]$   
    else  $C[h,j] := C[h+1,(j-1)/2] + B[h,j]$ 
```

Some More Issues

```
#pragma omp parallel for  
For i=1 to n do  
    A[i] = f(i);
```

- Load balancing
 - Consider a parallel for loop with the action being $f(i)$ for index i .

Some More Issues

```
#pragma omp parallel for  
For i=1 to n do  
    A[i] = f(i);
```

- Load balancing
 - Consider a parallel for loop with the action being $f(i)$ for index i .
 - It is likely that the time taken to finish execution $f(i)$ can vary for each i .
 - In that case, some threads finish their work early, and some others take longer.

Some More Issues

```
#pragma omp parallel for  
For i=1 to n do  
    A[i] = f(i);
```

- Load balancing
 - Consider a parallel for loop with the action being $f(i)$ for index i .
 - It is likely that the time taken to finish execution $f(i)$ can vary for each i .
 - In that case, some threads finish their work early, and some others take longer.
 - This is called as load imbalance.

Some More Issues

```
#pragma omp parallel for schedule (dynamic)  
For i=1 to n do  
    A[i] = f(i);
```

- OpenMP has additional features to help such settings.
- The modified program is as above.

Some More Issues

- There may be data dependencies across threads.
- For instance, at the same time that a thread is reading a shared variable, another thread may be writing to it.
- Or two threads may be writing at the same time to a shared variable.
- OpenMP has ways to handle this.

Some More Issues

- How to see the effect of parallel programs?
- One can measure the time taken on an execution.
- Highly dependent on the machine used, but presently no widely accepted abstract standard.
- Analytic models for parallel computing is an open research area.

Some More Issues

- Even practice is tricky.
- For small inputs, one may not see any benefit.
- This is due to some overheads involved.
- So, one should try for large enough inputs.
- For instance, the prefix sum program can be run on arrays of size 1 M and above.
- Typically one plots the runtime against input size.

One More Example

```
Program Merge(A, B)
Begin
  for i = 1 to n do in parallel
    RA[i] = BinSrch(B, A[i]) + i
  for i = 1 to n do in parallel
    RB[i] = BinSrch(A, B[i]) + i
  for i = 1 to n do in parallel
    C[RA[i]] = A[i]
  for i = 1 to n do in parallel
    C[RB[i]] = B[i];
end
```

- Recall our parallel version of the merge procedure.

One More Example

Program Merge(A, B)

Begin

#pragma omp parallel for shared(A,B) private(i)

for i = 1 to n do

RA[i] = BinSrch(B, A[i]) + i

#pragma omp parallel for shared(A,B) private(i)

for i = 1 to n do in parallel

RB[i] = BinSrch(A, B[i]) + 1

#pragma omp parallel for shared(A,B) private(i)

for i = 1 to n do in parallel

C[RA[i]] = A[i]

#pragma omp parallel for shared(A,B) private(i)

for i = 1 to n do in parallel

C[RB[i]] = B[i];

end

- Recall our parallel version of the merge procedure.

Summary

- Writing basic parallel programs is easy now.
- There are only a few constructs of openmp that we have to know.
- We will try a few in the tutorials and the laboratory sessions.
 - I tried the prefix sum and the merge program before class today.
- There is more but we may visit those issues when needed.
- Happy parallel programming 😊😊😊😊