

Breadth First Search

- Recall level order traversal of a tree.
 - Starting from the root, visits every vertex in a level by level manner.
- Let us develop breadth first search as an extension of level order traversal.
- A few questions to be answered before we develop breadth first search.

Breadth First Search

- **Question 1:** For a graph, no notion of a root vertex.
- So, where should BFS start from?

Breadth First Search

- **Question 1:** For a graph, no notion of a root vertex.
- So, where should BFS start from?
- So, have to specify a starting vertex. Typically denoted s .
- Still other problems exist.

Breadth First Search


- In a tree, using level order traversal, each vertex is visited also exactly once.
 - Why?

Breadth First Search

- In a tree, using level order traversal, each vertex is visited also exactly once.
 - Recall that a tree is connected and has no cycles.


Breadth First Search

- In a tree, using level order traversal, each vertex is visited also exactly once.
 - Recall that a tree is connected and has no cycles.
- In a graph, that is no longer guaranteed.
 - Start from $s = 2$ and do a level order traversal.




Breadth First Search

- In a tree, using level order traversal, each vertex is visited also exactly once.
 - Recall that a tree is connected and has no cycles.
- In a graph, that is no longer guaranteed.
 - Start from $s = 2$ and do a level order traversal
 - One of 1 or 5 visited more than once.




Breadth First Search

- **Question 2:** How to resolve that problem?



Breadth First Search

- **Question 2:** How to resolve that problem?
- Can remember if a vertex is already visited.
- Each vertex has a state among VISITED, NOT_VISITED, IN_PROGRESS.
- Why three states instead of just two?
 - Need them for a later use.




Breadth First Search

- **Question 3:** Can all vertices be reached from s ?

Breadth First Search

- Question 3: Can all vertices be reached from s ?
- For example, when $s = 2$, vertex 3 can never be visited.
- What to do with those vertices?
- Answer depends on the idea behind graph searching via BFS.



Breadth First Search

- The basic idea of breadth first search is to find the least number of edges between s and any other vertex in G .
 - The same property holds for level order traversal of a tree also with s as the root.
- Starting from s , we can thus visit vertices of distance k before visiting any vertex of distance $k+1$.
- For that purpose, define $d_s(v)$ to be the least number of edges between s and v in G .

Breadth First Search


- So, for vertices v that are not reachable from s , can say that $d_s(v)$ is ∞
- Alike a level order traversal of a tree, can use a queue to store vertices in progress.

BFS Procedure

```
Procedure BFS(G)
for each  $v \in V$  do
 $\pi(v) = \text{NIL}$ ;  $\text{state}[v] = \text{NOT\_VISITED}$ ;  $d(v) = \infty$ ;
End-for
 $d[s] = 0$ ;  $\text{state}[s] = \text{IN\_PROGRESS}$ ;  $\pi[s] = \text{NIL}$ ,
 $Q = \text{EMPTY}$ ;  $Q.\text{Enqueue}(s)$ ;
While  $Q$  is not empty do
 $v = Q.\text{Dequeue}()$ ;
for each neighbour  $w$  of  $v$  do
    if  $\text{state}[w] = \text{NOT\_VISITED}$  then
         $\text{state}[w] = \text{IN\_PROGRESS}$ ;  $\pi[w] = v$ ;
         $d[w] = d[v] + 1$ ;  $Q.\text{Enqueue}(w)$ ;
    end-if
end-for
 $\text{state}[v] = \text{FINISHED}$ 
end-while
```

BFS Example

- Start from $s = 2$.



	1	2	3	4	5	6
d :	∞	0	∞	∞	∞	∞
π :	—	—	—	—	—	—

BFS – Additional Details

- What is the runtime of BFS?


BFS – Additional Details

- What is the runtime of BFS?
 - How many times does each vertex enter the queue?
 - Each edge is considered only once.
- Therefore, the runtime of BFS should be $O(m + n)$.

BFS – Additional Details

- The π value of a vertex v denotes the vertex u that **discovered** v .
- The π values maintained during BFS can be used to define a subgraph of G as follows.
- Define the predecessor subgraph of $G = (V, E)$ as
 - $G_\pi = (V_\pi, E_\pi)$ where
 - $V_\pi = \{v \in V : \pi(v) \neq \text{NULL}\} \cup \{s\}$, i.e., all vertices reached during a BFS from s , and
 - $E_\pi = \{(\pi(v), v) \in E : v \in V_\pi - \{s\}\}$, directed edges from the parent of a vertex to the vertex.

BFS Example Contd...



Properties of BFS


- Consider the time at which a vertex v has entered the queue.
- The state of v at that instant changes from NOT_VISITED to IN_PROGRESS.
- $d_s(v)$ changes to a finite value, and
- $d_s(v)$ can never change after that instant.

Classifying Edges

- Can classify edges of G according to BFS from a given s as follows.
- The edges of E_π are also called as **tree edges**.
- It holds that for a tree edge (u, v) , $d(v) = d(u) + 1$.
- The edges of $E_N := E \setminus E_\pi$ are called as **non-tree edges**.
- These edges can be further classified as follows.

Classifying Edges


- Identify the tree- and the non-tree edges according to a BFS on the following graph. Choose vertex number 3 as the start vertex.




- Pick vertices in their order.

Classifying Edges – The Non-Tree Edges

- First, consider the predecessor subgraph. It is a tree. Call this tree as T_{BFS} .
- Tree edges according to BFS share a parent-child relationship.
- For any pair of vertices u, v :
 - Either they share an ancestor-descendant relation in T_{BFS} .
 - Or they do not.




Classifying Edges – The Non-Tree Edges



- For any pair of vertices u, v :
 - Either they share an ancestor-descendant relation in T_{BFS} .
 - Or they do not.
 - (u, v) called as a **cross edge**. Examples (d, i) and (b, a) .

Classifying Edges – The Non-Tree Edges




- For any pair of vertices u, v with (u,v) an edge in G :
 - Either they share an ancestor-descendant relation in T_{BFS} .
 - If u is an ancestor of v , then (u,v) is a **forward edge**.
 - If u is a descendant of v , then (u,v) is a **back edge**.

Directed or Undirected

- Most of the above observations hold even if G is directed.
 - The classification in fact makes more sense for directed graphs.
 - There can be back edges, but no forward edges.
- Can thus extend the notion of BFS to directed graphs.

Complete Example

- Perform BFS on the directed graph below with vertex a as the start vertex.
- Classify the edges of the graph according to the BFS.



BFS – Colors instead of States

- It is common to associate colors to the three states.
 - GREEN : Done vertices, VISITED
 - ORANGE : In progress/ In Queue
 - RED: Not visited yet.


Towards Weighted BFS

- So, far we have measured $d_s(v)$ in terms of number of edges in the path from s to v .
- Equivalent to assuming that each edge in the graph has equal (unit) weight.
- But, several settings exist where edges may have unequal weights.

Towards Weighted BFS

- Consider a road network.
- Junctions can be vertices and roads can be edges.
- Can use such a graph to find the **best** way to reach from point A to point B.
- **Best** here can mean shortest distance/shortest delay/....
- Clearly, all edges need not have the same distance/delay/.

Towards Weighted BFS



A Few Problems

- Problem I : Given two points u and v , find the shortest distance between them.
- Problem II : Given a starting point s , find the shortest distance from s to all other points.
- Problem III : Find the shortest distance between all pairs of points.

A Few Problems

- Turns out that Problem I is not any easier than Problem II.
- Problem III is definitely harder than Problem II.
- We shall study problem II, and possibly Problem III.

Weighted Graphs

- The setting is more general.
- A weighted graph $G = (V, E, W)$ is a graph with a weight function $W : E \rightarrow \mathbb{R}$.
- Weighted graphs occur in several settings
 - Road networks
 - Internet

Problem II : Single Source Shortest Paths

- Problem II is also called the single source shortest paths problem.
- Let us extend BFS to solve this problem.
- Notice that BFS solves the problem when all the edge weights are 1.
 - Hence the reason to extend BFS

SSSP

- Extensions needed
 - 1. Weights on edges


SSSP

- Extensions needed
 1. Weights on edges
 2. How to know when a node is finished.

SSSP

- Extensions needed
 - 1. Weights on edges
 - 2. How to know when a node is finished.
- For a vertex v , $d_s(v)$ will now refer to the shortest distance from s to v .
- Initially, like in BFS, $d_s(v) = \infty$ for all vertices v except s , and $d_s(s) = 0$.

Weighted BFS




- Update $d_s(v)$ with weights.
- Also, weights on edges mean that if v is a neighbor of w in the shortest path from s to w , then $d_s(w) = d_s(v) + W(v,w)$.
 - Instead of $d_s(w) = d_s(v) + 1$ as in BFS.
- We will call this as the **first change to BFS**.

SSSP

- Notice that in BFS a node has three states :
NOT_VISITED, VISITED, IN_QUEUE
- A vertex in VISITED state should have no more changes to $d_s()$ value.
- What about a vertex in IN_QUEUE state?
 - such a vertex has some finite value for $d_s(v)$.
 - Can $d_s(v)$ change for such vertices?
 - Consider an example.


Weighted BFS

- Consider $s = 2$ and perform weighted BFS.



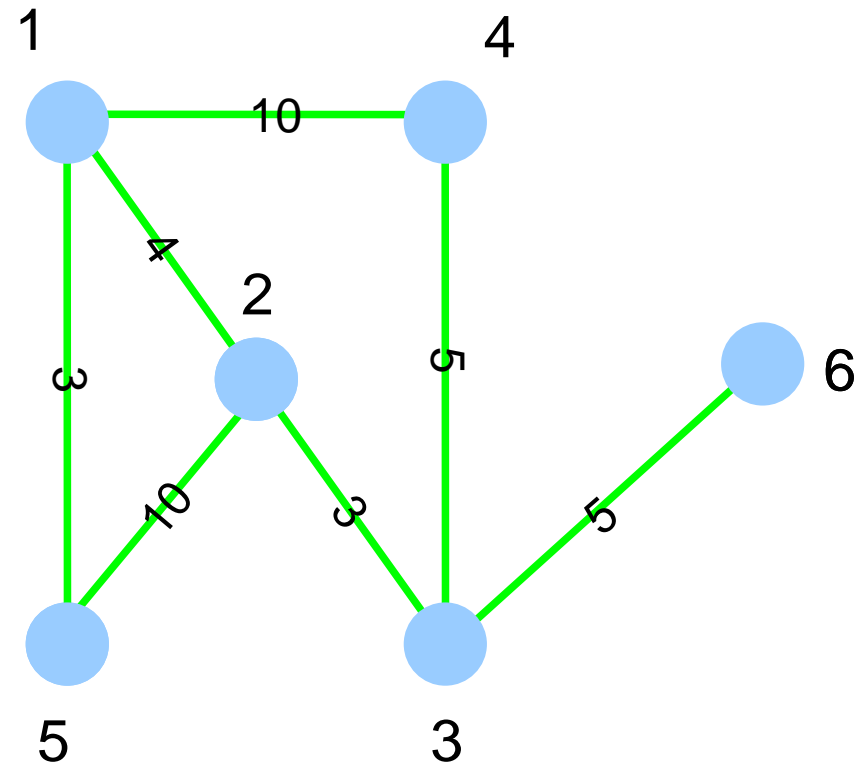
Weighted BFS

- Consider $s = 2$.
- From s , we will enqueue 1, 5, and 3 with $d(1) = 4$, $d(5) = 10$, $d(3) = 3$, in that order.
- While vertex 5 is still in queue, can visit 5 from vertex 1 also.



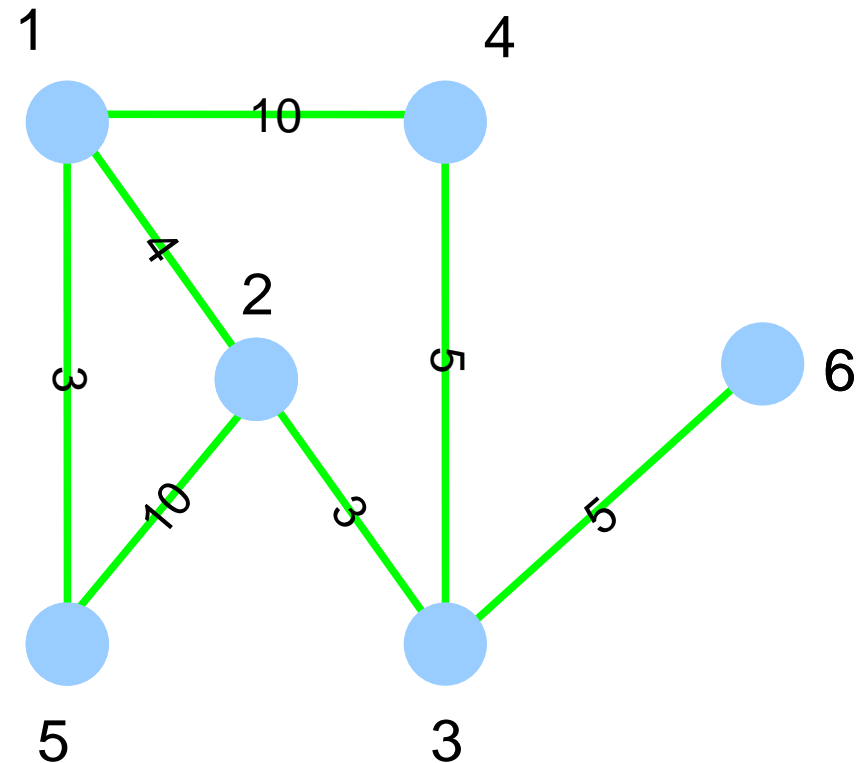
Weighted BFS

- Moreover, the weight of the edge 2- 5 is 10 whereas there is a shorter path from 2 to 5 via the path 2 – 1 – 5.
- So, it suggests that $d(v)$ should be changed while v is still in the queue.



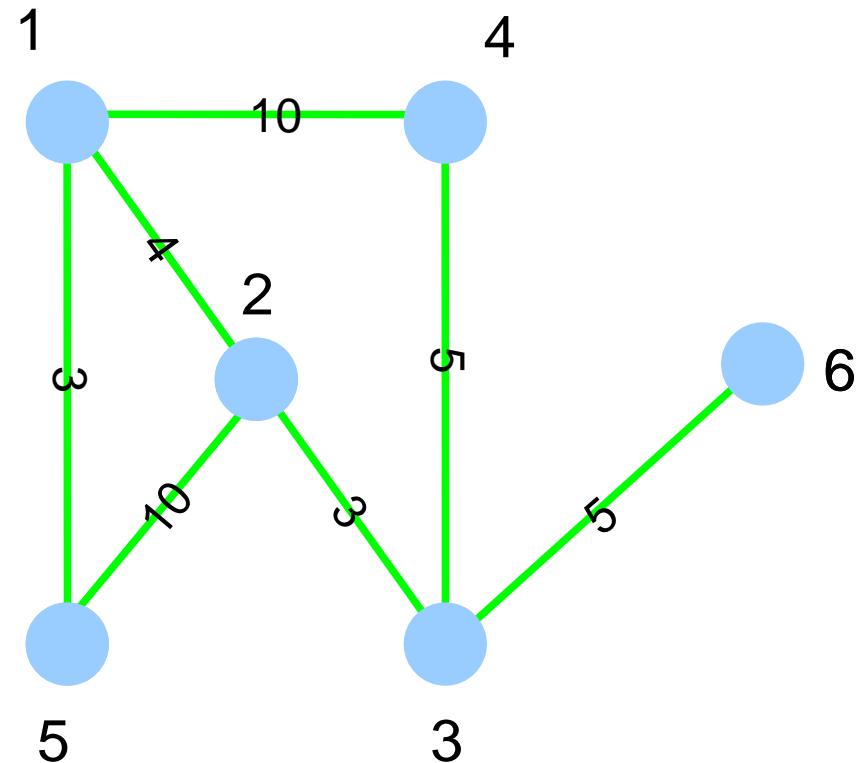
Weighted BFS

- Update $d(v)$ for v in queue also.
- While v is in queue, we can check if $d(v)$ is more than the distance along the new path.
- If so, then update $d(v)$ to the new smaller value.
- Change 2 to BFS.



Weighted BFS

- Does that suffice?
- In the same example, if we change the order of vertices from 1, 5, 3 to 5, 1, 3, then vertex 5 will not be in queue when 1 is removed from the queue.



Weighted BFS

- So, the simple fix to change $d(v)$ while v is still in queue does not work.
- May need to update $d(v)$ even when v is not in queue?
 - But how long should we do so?

Weighted BFS

- Can do so as long as there are changes to some $d(v)$?
 - No need of a queue then, in this case really.
- Will this ever stop?
- Indeed it does. Why?
 - Intuitively, there are only a finite number of edges in any shortest path.

Weighted BFS

- Why does this ever stop?
- Consider a vertex v and the path from s to v of the least cost.

An Algorithm for SSSP


```
Algorithm SSSP(G,s)
begin
  for all vertices v do
     $d(v) = \infty$ ;  $\pi(v) = \text{NIL}$ ;
  end-for
   $d(s) = 0$ ;
  for n-1 iterations do
    for each edge (v,w) do
      if  $d(w) > d(v) + W(v,w)$  then
         $d(w) = d(v) + W(v,w)$ ;  $\pi(w) = v$ ;
      end-if
    end-for
  end-for
end
```

Algorithm SSSP

- The above algorithm is called the Bellman-Ford algorithm.
- The algorithm requires $O(mn)$ time.
 - For each of the $n-1$ iterations, we consider each edge once.
 - Has $O(1)$ compute per edge.
- Just as in BFS, works also on directed graphs.
- Forms the basis of several algorithms for the Internet.

Example Algorithm SSSP

- Start vertex = d. Employ the Bellman-Ford algorithm to find shortest path from d to all other vertices.




Thinking about the Bellman-Ford Algorithm

```
Algorithm SSSP(G,s)
begin
  for all vertices v do  $d(v) = \infty$ ;  $\pi(v) = \text{NIL}$ ;
   $d(s) = 0$ ;
  for n-1 iterations do
    for each edge (v,w) do
      if  $d(w) > d(v) + W(v,w)$  then
         $d(w) = d(v) + W(v,w)$ ;  $\pi(w) = v$ ;
end
```

- Why n-1 iterations are required?
- Let us prove the following via induction.

Thinking about the Bellman-Ford Algorithm


- Consider the source vertex s .
- For s , $d(s) = 0$ is the best possible result.
- So, s is FINISHED.



- Now consider a vertex v such that the shortest path from s to v contains only one edge, say (s,v) .
- The edge (s,v) appears at some iteration of the second for loop in the first iteration of the main loop.
- At that point, $d(v)$ is set correctly.

Thinking about the Bellman-Ford Algorithm

- Consider the source vertex s .
- For s , $d(s) = 0$ is the best possible result.
- So, s is FINISHED.




- Now consider a vertex v such that the shortest path from s to v contains only one edge, say (s,v) .
- The edge (s,v) appears at some iteration of the second for loop in the first iteration of the main loop.
- At that point, $d(v)$ is set correctly.
- Does that mean that all neighbors of s FINISH in one iteration?

Thinking about the Bellman-Ford Algorithm

- In that fashion, let every vertex v with a shortest path having at most i edges enter the FINISHED state at the end of i iterations.
- This certainly holds for $i = 0$. (and $i = 1$ too!)
- Can we use induction to continue the proof?

The Proof

- In pictures...



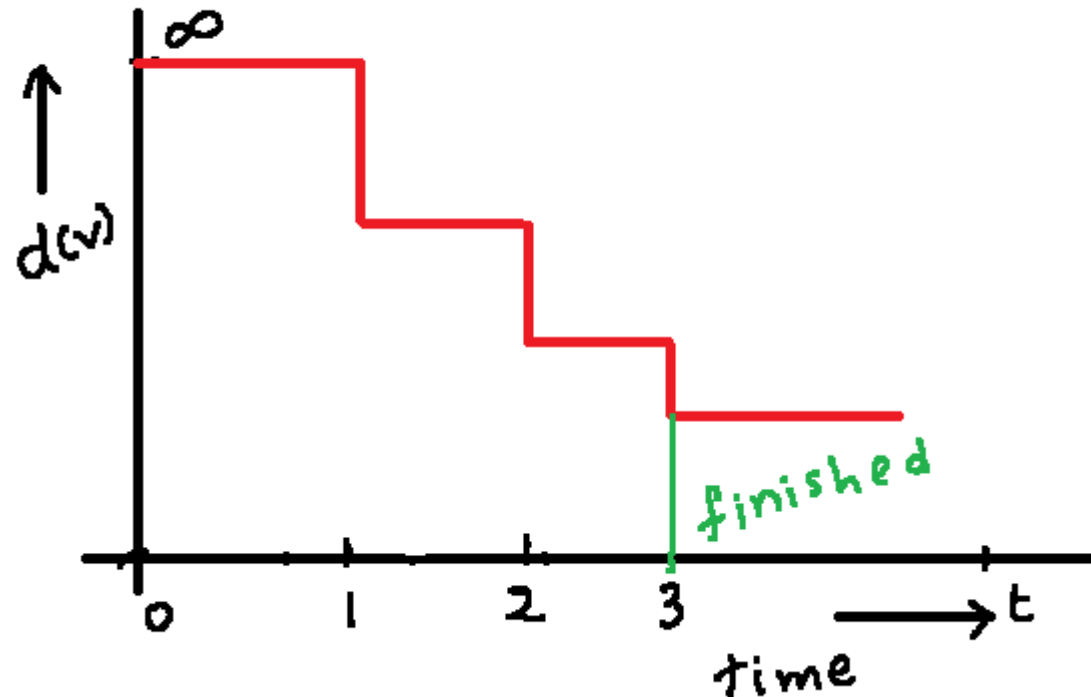
Algorithm SSSP

- The time taken by the Bellman-Ford algorithm is too high compared to that of BFS.
- Can we improve on the time requirement?
- Most of the time is due to

Algorithm SSSP


- The time taken by the Bellman-Ford algorithm is too high compared to that of BFS.
- Can we improve on the time requirement?
- Most of the time is due to
 - Repeatedly considering edges, and as a result
 - Updating $d(v)$ possibly many times
- Need to know how to stop updating $d(v)$ for any vertex v .
- This is what we will develop next.

To Improve the Runtime



- When is a vertex FINISHED?
- When no further shorter path can be found to v from s .
 - Equivalently, when $d(v)$ can no longer decrease.

A Considered Edge



```

void process(e) /*e = (v,w)*/
begin
    if d(w) > d(v) + W(v,w) then
        d(w) = d(v) + W(v,w)
    end
end
  
```

- We say an edge $e = (v, w)$ is **considered** if the above routine is executed for e .
- The impact is to possibly lower $d(w)$, indicating that a **better** path to w from s is available via v .


To Improve the Runtime

- For this to happen, consider the following.
 - A few vertices, say S , are FINISHED.
 - Plus, all the edges with at least one endpoint in S are the only edges considered.
 - Other vertices in $V \setminus S$, have some $d()$ value.

To Improve the Runtime


- For this to happen, consider the following.
 - Let each edge have a positive weight.
 - A few vertices, say S , are FINISHED.
 - Plus, all the edges with at least one endpoint in S are the only edges considered.
 - Other vertices in $V \setminus S$, have some $d()$ value.
 - Which of these cannot improve $d()$ any more?

The Setting




- **Green** vertices are FINISHED.
- **Red** edges, edges with at least one end point as a green vertex are the ONLY edges PROCESSED.
- Numbers on black vertices indicate their $d()$ value using only green vertices as intermediate vertices.

The Setting



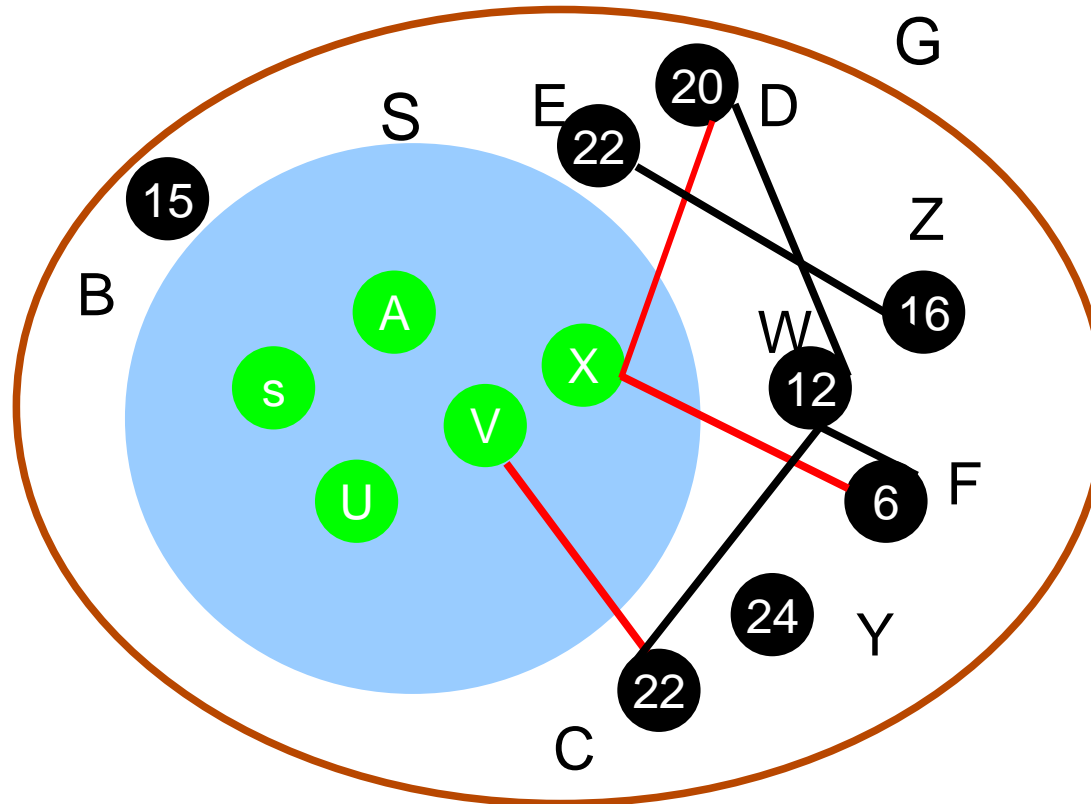
- Suppose we want to add one more vertex to the set S.
- Which of the black vertices is FINISHED?

The Setting




- Notice that there could be edges between the black vertices also.
 - None of them are processed so far.

The Setting




- Consider the vertex v with the smallest $d()$ value among the black vertices.
- Any more decrease to $d(v)$ would involve using at least one more edge between two black vertices.

The Setting



- Consider the vertex v with the smallest $d()$ value among the black vertices.
- Any more decrease to $d(v)$ would involve using at least one more edge between two black vertices.
- But all edge weights are positive.

The Setting



- Therefore, such a vertex with the smallest $d()$ value among the black vertices can no longer decrease its $d()$ value.