

Our First Data Structure

- The story so far
 - We understand the need for data structures
 - We have seen a few analysis techniques
- This week we will
 - Attempt a definition of what is a data structure
 - See a very simple data structure, and
 - Advanced applications of this data structure.

A Data Structure

- How should we view a data structure?
- From an implementation point of view
 - Should implement a set of operations
 - Should provide a way to store the data in some form.
- From a user point of view
 - should use it as a black box
 - call the operations provided
- Analogy : C programming language has a few built-in data types.
 - int, char, float, etc.

Analogy

- Consider a built-in data type such as int
- A programmer
 - can store an integer within certain limits
 - can access the value of the stored integer
 - can do other operations such add, subtract, multiply, ...
- Who is implementing the data structure?
 - A compiler writer.
 - Interacts with the system and manages to store the integer.


An Abstract Data Type

- A data structure can thus be looked as an **abstract data type**.
- An abstract data type specifies a set of operations for a user.
- The implementor of the abstract data type supports the operations in a suitable manner.
- One can thus see the built-in types also as abstract data types.

The Array as a Data Structure

- Suppose you wish to store a collection of like items.
 - say a collection of integers
- Will access any item of the collection.
- May or may not know the number of items in the collection.
- Example settings:
 - store the scores on an exam for a set of students.
 - store the temperature of a city over several days.

The Array as a Data Structure



- In such settings, one can use an array.
- An array is a collection of like objects.
 - Usually denoted by upper case letters such as A, B.
- Let us now define this more formally.

The Array ADT

- Typical operations on an array are:
 - `create(name, size)` : to create an array data structure,
 - `ElementAt(index, name)` : to access the element at a given index i
 - `size(name)` : to find the size of the array, and
 - `print(name)` : to print the contents of the array
- Note that in most languages, `elementAt(index, name)` is given as the operation `name[index]`.

The Array Implementation

```
Algorithm Create(int size, string name)
begin
  name = malloc(size*sizeof(int));
end
```

```
Algorithm ElementAt(int index,
  string name)
begin
  return name[i];
end
```

```
Algorithm Print(string name)
begin
  for i = 1 to n do
    printf("%d t",
      name[i]);
  end-for;
end;
```

```
Algorithm size(string name)
begin
  return size;
end;
```


Further Operations

- The above operations are quite fundamental.
- Need further operations for most problems.
- We will now see some such operations.

Sorting

- Sorting is a fundamental concept in Computer Science.
 - several application and a lot of literature.
 - We shall see an algorithm for sorting.

QuickSort

- The quick sort algorithm designed by Hoare is a simple yet highly efficient algorithm.
- It works as follows:
 - Start with the given array A of n elements.
 - Consider a pivot, say $A[1]$.
 - Now, partition the elements of A into two arrays A_L and A_R such that:
 - the elements in A_L are less than $A[1]$
 - the elements in A_R are greater than $A[1]$.
 - Sort A_L and A_R , recursively.

How to Partition?

- Here is an idea.
 - Suppose we take each element, compare it with $A[1]$ and then move it to A_L or A_R accordingly.
 - Works in $O(n)$ time.
 - Can write the program easily.
 - But, recall that space is also an resource. The above approach requires extra space for the arrays A_L and A_R
 - A better approach exists.

Algorithm Partition

```
Procedure Partition(A,n)
begin
  pivot = A(n);
  less = 0; more = 1;
  for more = 1 to n do
    if A(more) < pivot then
      less++;
      swap(A(more), A(less));
    end
  swap A[less+1] with A[n];
end
```

- Algorithm Partition is given above.

Example

15	24	37	12	18	11	42	25	22
----	----	----	----	----	----	----	----	----

15	24	37	12	18	11	42	25	22
----	----	----	----	----	----	----	----	----

15	24	37	12	18	11	42	25	22
----	----	----	----	----	----	----	----	----

15	24	37	12	18	11	42	25	22
----	----	----	----	----	----	----	----	----

15	12	37	24	18	11	42	25	22
----	----	----	----	----	----	----	----	----

15	12	18	24	37	11	42	25	22
----	----	----	----	----	----	----	----	----

15	12	18	11	37	24	42	25	22
----	----	----	----	----	----	----	----	----

15	12	18	11	37	24	42	25	22
----	----	----	----	----	----	----	----	----

15	12	18	11	37	24	42	25	22
----	----	----	----	----	----	----	----	----

Practice Problem


Partition the elements below around the last element as the pivot.

Show all your work.

24, 41, 9, 18, 36, 16, 19, 48, 20

Correctness by Loop Invariants

- Consider the Partition algorithm as an example.
- The Partition algorithm partitions the input data set into three parts around a pivot.



Correctness by Loop Invariants

- How to prove that the above algorithm is correct
 - We shall use a loop invariant.
- How do we come up with a loop invariant?
 - Study the loop for its purpose and construction.
 - What property of the loop can we seek during every iteration?
 - Formalize such a set of statements.

Correctness by Loop Invariants

- Statement of the loop invariant


After k iterations of the loop, the following hold:

- Elements $A(1)$ to $A(\text{left})$ are less than the pivot.
- Elements $A(\text{left}+1)$ to $A(\text{right})$ are greater than the pivot.
- Elements $A(\text{right}+1)$ to $A(n-1)$ are not classified.
- $A(n) = \text{pivot}$.

Using a Loop Invariant (LI)

- We show three things with respect to a loop invariant.
- **Initialization**: The LI is true prior to the first iteration of the loop.
- **Maintenance**: If the LI holds true before a particular iteration then it is true before the start of the next iteration.
- **Termination**: Upon termination of the loop, the LI can be used to state some property of the algorithm and establish its correctness.

The Basic Step in Partition




Correctness by Loop Invariants

- Initialization: Check that the loop invariant true before the start of the loop.
- In our example, $\text{left} = 0$, $\text{right} = 0$ before the start of the loop.
- So, the four conditions are met.

Correctness by Loop Invariants


- Maintenance: Assume that the loop invariant is true for the past k iterations.
- Alike induction step, we need to show that also for the $k+1$ iterations, the loop invariant holds.
- Consider the actions in the loop and their effect on the loop invariant.

The Basic Step in Partition



- The main action in the loop is the comparison of $A[k+1]$ with $A[n]$.
- Consider the case when $A[k+1] < A[n]$.

The Basic Step in Partition



- Consider the case when $A[k+1] > A[n]$

Practice Problem:: Loop Invariants

- Consider the following algorithm. What does it do? Formulate an appropriate loop invariant, and show that the algorithm is correct.

Algorithm WhatIsThis(X)

Begin

 int i = 1;

 while (i <= n)

 int j = i+1;

 while (j <= n)

 if (X[i] > X[j])

 y = X[i]; X[i] = X[j]; X[j] = y;

 j++;

 i++;

End

Practice Problem:: Loop Invariants

- The algorithm is sorting X . The procedure is called bubble sort.
- One possible loop invariant:
 - For the first while-loop: At the end of i iterations, the elements of X are such that $X[1] \leq X[2] \leq \dots X[i]$.
 - For the second while-loop: At the end of k iterations of the loop for $k = i+1, i+2, \dots, n$, the elements $X[i+1], X[i+2], \dots, X[k]$ are all smaller than $X[i]$.

Analyzing Quick Sort

- Suppose we run quick sort with $A[n]$ as the pivot.
- Let A_L and A_R be the two subarrays obtained after partitioning.
- What is the time taken by quicksort?
- As a recurrence relation, $T(n) = T(|A_L|) + T(|A_R|) + O(n)$.
- To be able to solve this recurrence relation, need to know the sizes of arrays A_L and A_R .

Analyzing Quick Sort

- We know that $|A_L| + |A_R| = n-1$.
- But, if the pivot is such that all elements are smaller (or larger) than the pivot, then $|A_L|$ (or $|A_R|$) = $n-1$.
- The recurrence relation in that case is
$$T(n) = T(n-1) + O(n).$$
- Suppose the same situation happens over every recursive call. So, the above recurrence relation holds during every recursive call.

Example Bad Case

12 15 17 24 29 36 | 42
 A_L

12 15 17 24 29 | 36
 A_L

12 15 17 24 | 29
 A_L

12 15 17 | 24
 A_L

12 15 | 17
 A_L

Analysis of Quick Sort

- Find the solution to the recurrence relation

$$T(n) = T(n-1) + O(n)$$

Analysis of Quick Sort

- Is it always that bad?
- What if the pivot is such that each recursive iteration, the sizes of $|A_L|$ and $|A_R|$ is exactly the same?
- The recurrence relation then stands as:
$$T(n) = 2T(n/2) + O(n).$$
- Solve this recurrence relation.

Analysis of Quick Sort

- Which element as the pivot ensures that the sizes of $|A_L|$ and $|A_R|$ are exactly the same?
- Can that happen in every run?

Analysis of Quick Sort

- Which element as the pivot ensures that the sizes of $|A_L|$ and $|A_R|$ are exactly the same?
- Can that happen in every run?
- In general, if the sizes of $|A_L|$ and $|A_R|$ are such that they are a constant away from each other, then the recurrence relation is:

$$T(n) = T(an) + T((1-a)n) + O(n)$$

where a is a constant < 1 .

- Can you solve this recurrence relation?

Analysis of Quick Sort

- In practice, it turns out that most often the partitions are not too skewed.
- So, quick sort runs in $O(n \log n)$ time almost always.

Another Operation – Prefix Sums

- Consider any associative binary operator, such as $+$, and an array A of elements over which o is applicable.
- The prefix operation requires us to compute the array S so that $S[i] = A[1] + A[2] + \dots + A[i]$.
- The prefix operation is very easy to perform in the standard sequential setting.

Sequential Algorithm for Prefix Sum

```
Algorithm PrefixSum(A)
```

```
  S[1] = A[1];
```

```
  for i = 2 to n do
```

```
    S[i] = A[i] + S[i-1]
```

```
  end-for
```

- Example $A = (3, -1, 0, 2, 4, 5)$
- $S[1] = 3$.
- $S[2] = -1 + 3 = 2$, $S[3] = 0 + 2 = 2, \dots$
- The time taken for this program is $O(n)$.

Our Interest in Prefix

- The world is moving towards parallel computing.
- This is necessitated by the fact that the present sequential computers cannot meet the computing needs of the current applications.
- Already, parallel computers are available with the name of multi-core architectures.
 - Majority of PCs today are at least dual core.

Our Interest in Prefix

- Programming and software has to wake up to this reality and have a rethink on the programming solutions.
- The parallel algorithms community has fortunately given a lot of parallel algorithm design techniques and also studied the limits of parallelizability.
- How to understand parallelism in computation?

Parallelism in Computation

- Think of the sequential computer as a machine that executes jobs or instructions.
- With more than one processor, can execute more than one job (instruction) at the same time.
 - Cannot however execute instructions that are dependent on each other.

Parallelism in Computation

- This opens up a new world where computations have to specified in parallel.
- Sometimes have to rethink on known sequential approaches.
- Prefix computation is one such example.
 - Turns out that prefix sum is a fundamental computation in the parallel setting.
 - Applications span several areas.

Parallelism in Computation

- The obvious sequential algorithm for prefix sums does not have enough independent operations to benefit from parallel execution.
- Computing $S[i]$ requires computation of $S[i-1]$ to be completed.
- Have to completely rethink for a new approach.

Parallel Prefix

- Consider the array A and produce the array B of size $n/2$ where $B[i] = A[2i - 1] + A[2i]$.
- Imagine that we recursively compute the prefix output wrt B and call the output array as C .
- Thus, $C[i] = B[1] + B[2] + \dots + B[i]$. Let us now build the array S using the array C .

Parallel Prefix

- For this, notice that for even indices i , $C[i] = B[1] + B[2] + \dots + B[i] = A[1] + A[2] + \dots + A[2i]$, which is what $S[2i]$ is to be.
- Therefore, for even indices of S , we can simply use the values in array C .

Parallel Prefix

- For this, notice that for even indices i , $C[i] = B[1] + B[2] + \dots + B[i] = A[1] + A[2] + \dots + A[2i]$, which is what $S[2i]$ is to be.
- Therefore, for even indices of S , we can simply use the values in array C .
- The above also suggests that for odd indices of S , we can apply the $+$ operation to a value in C and a value in A .

Parallel Prefix Example

- $A = (3, 0, -1, 2, 8, 4, 1, 7)$
- $B = (3, 1, 12, 8)$
 - $B[1] = A[1] + A[2] = 3 + 0 = 3$
 - $B[2] = A[3] + A[4] = -1 + 2 = 1$
 - $B[3] = A[5] + A[6] = 8 + 4 = 12$
 - $B[4] = A[7] + A[8] = 1 + 7 = 8$
- Let C be the prefix sum array of B , computed recursively as $C = (3, 4, 16, 24)$.
- Now we use C to build S as follows.

Parallel Prefix Example

- $S[1] = A[1]$, always.
- $C[1] = B[1] = A[1] + A[2] = S[2]$
- $C[2] = B[1] + B[2] = A[1] + A[2] + A[3] + A[4] = S[4]$
- $C[3] = B[1] + B[2] + B[3]$
 $= A[1] + A[2] + A[3] + A[4] + A[5] + A[6] = S[6]$
- That completes the case for even indices of S .
- Now, let us see the odd indices of S .

Parallel Prefix Example

- Consider, $S[3] = A[1] + A[2] + A[3]$
$$= (A[1] + A[2]) + A[3]$$
$$= S[2] + A[3].$$
- Similarly, $S[5] = S[4] + A[5]$ and $S[7] = S[6] + A[7]$.
- Notice that if $C[2]$, $C[4]$, and $C[6]$ are known, the computation at odd indices is independent for every odd index.

Parallel Prefix Algorithm

Algorithm Prefix(A)

begin

Phase I: Set up a recursive problem

for $i = 1$ to $n/2$ do

$B[i] = A[2i - 1] \circ A[2i];$

end-for

Phase II: Solve the recursive problem

Solve Prefix(B) into C;

Phase III: Solve the original problem

for $i = 1$ to n do

if $i = 1$ then $S[1] = A[1];$

else if i is even then $S[i] = C[i/2];$

else if i is odd then $S[i] = C[(i - 1)/2] \circ A[i];$

end-for

end

Analyzing the Parallel Algorithm

- Can use the asymptotic model developed.
- Identify which operations are independent.
- These all can be done at the same time provided resources exist.
- In our algorithm
 - Phase I : has $n/2$ independent additions.
 - Phase II : using our knowledge on recurrence relations, this takes time $T(n/2)$.
 - Phase III : Here, we have another n independent operations.

Analyzing the Parallel Algorithm

- How many independent operations can be done at a time?
 - Depends on the number of processors available.
- Assume that as many as n processors are available.
- Hence, phase I can be done in $O(1)$ time totally.
- Phase II can be done in time $T(n/2)$
- Phase III can be done in $O(1)$ time.

Analyzing the Parallel Algorithm

- Using the above, we have that
 - $T(n) = T(n/2) + O(1)$
 - Using Master's theorem, can also see that the solution to this recurrence relation is $T(n) = O(\log n)$.
- Compared to the sequential algorithm, the time taken is now only $O(\log n)$, when n processors are available.

How Realistic is Parallel Computation?

- Our analysis suggests that the computation takes only $O(\log n)$ time, but we need n processors for this.
- Cannot ensure that the number of processors also grow with the input size.
- In practice, the number of processors on a machine does not change!

How Realistic is Parallel Computing

- The idea of the parallel algorithm is to show the extent of parallelism available in the computation.
- Plus, if there are fewer processors than what is required, can always simulate more processors.
- For instance, if there are p processors and n processors are required, then each of the p processors simulates the actions of n/p processors.

How Realistic is Parallel Computation

- Practical experience indicates that this is a viable proposition.

Merge Sort and Parallel Merge Sort

- Another sorting technique.
- Based on the divide and conquer principle.
- We will first explain the principle and then apply it to merge sort.


Divide and Conquer

- Divide the problem P into $k \geq 2$ sub-problems P_1, P_2, \dots, P_k .
- Solve the sub-problems P_1, P_2, \dots, P_k .
- Combine the solutions of the sub-problems to arrive at a solution to P .

Basic Techniques – Divide and Conquer

- A useful paradigm with several applications.
- Examples include merge sort, convex hull, median finding, matrix multiplication, and others.
- Typically, the sub-problems are solved recursively.
 - Recurrence relation

$$T(n) = D(n) + \sum_i T(n_i) + C(n)$$



Divide and Conquer

- Combination procedure : Merge

→ 8 10 12 27 → 15 17 24 32

8

8 10

8 10 12

8 10 12 15

8 10 12 15 17

8 10 12 15 17 24

8 10 12 15 17 24 27

8 10 12 15 17 24 27 32

Algorithm Merge

Algorithm Merge(L, R)

// L and R are two sorted arrays of size n each.

// The output is written to an array A of size 2n.

int i=1, j=1;

L[n+1] = R[n+1] = MAXINT; // so that index does not
// fall over

for k = 1 to 2n do

 if L[i] < R[j] then

 A[k] = L[i]; i++;

 else A[k] = R[j]; j++;

end-for

Algorithm Merge – Practice Problem

- Analyze the merge algorithm for its runtime.

Algorithm Merge – Practice Problem

- Analyze the merge algorithm for its runtime.
- Notice that there is a for-loop of $2n$ iterations.
- The number of comparisons performed is $O(n)$.
- Hence, the total time is $O(n)$.
- Is it correct?

Correctness of Merge

- We can argue that the algorithm Merge is correct by using the following **loop invariant**.
- At the beginning of every iteration
 - $L[i]$ and $R[j]$ are the smallest elements of L and R respectively that are not copied to A .
 - $A[1..k - 1]$ is in sorted order and contains the smallest $i - 1$ and $j - 1$ elements of L and R respectively.
- Need to verify these statements.

Correctness of Merge

- **Initialization** : At the start we have $i = j = 1$ and A is empty. So both the statements of the LI are valid.
- **Maintenance** : Let us look at any typical iteration k .
Let $L[i] < R[j]$.
- By induction, these are the smallest elements of L and R respectively and are not put into A .
- Since we add $L[i]$ to A at position k and do not advance j the two statements of the LI stay valid after the completion of this iteration.

Correctness of Merge

- **Termination** : At termination $k = l + r + 1$ and by the second statement we have that A contains $k - 1 = l + r$ elements of L and R in sorted order.
- Hence, the algorithm Merge correctly merges the sorted arrays L and R .

From Merging to Sorting

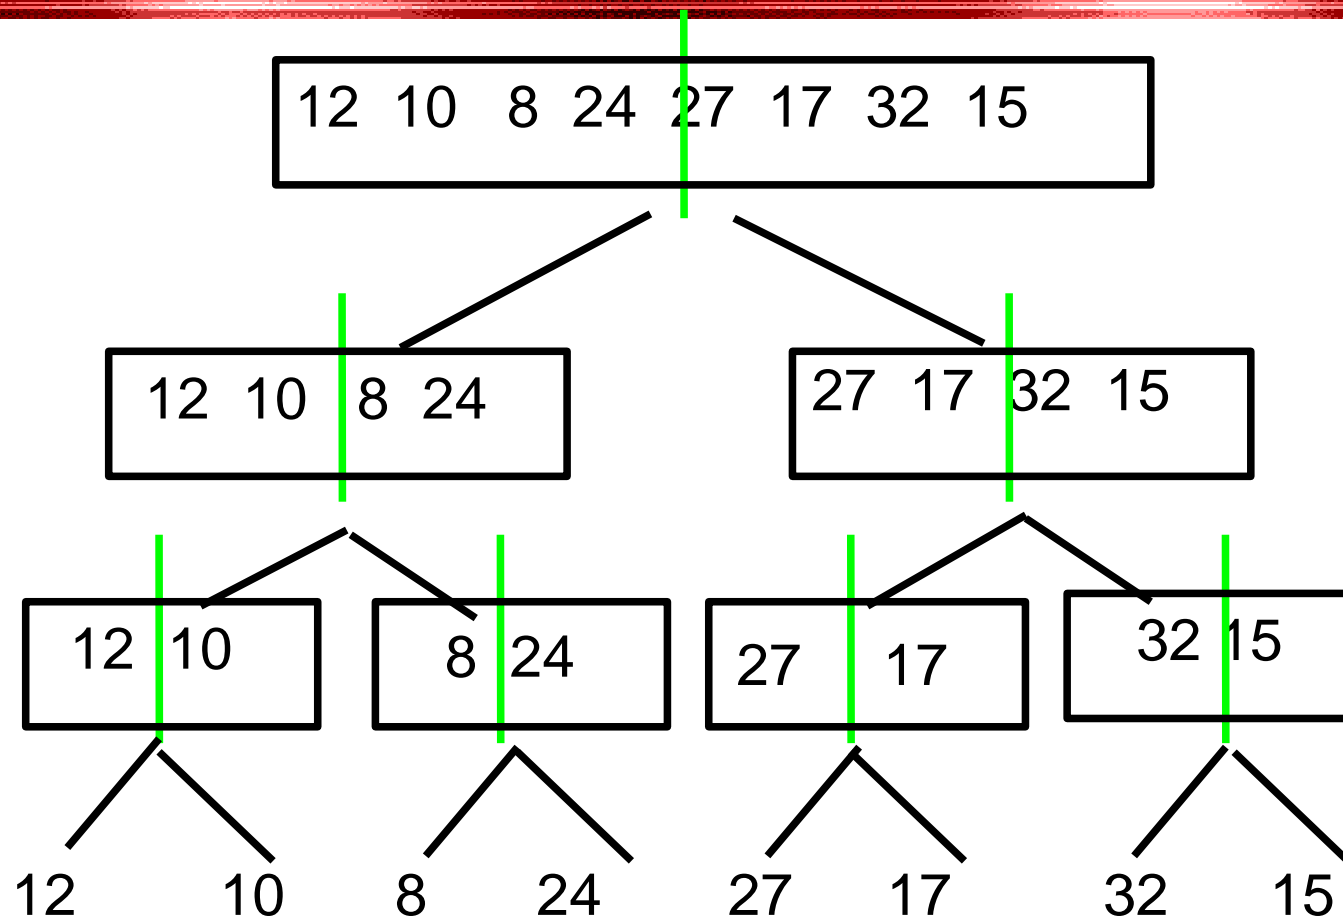
- How to use merging to finally sort?
- Using the divide and conquer principle
 - Divide the input array into two halves.
 - Sort each of them.
 - Merge the two sub-arrays. This is indeed procedure Merge.
- The algorithm can now be given as follows.

Algorithm MergeSort

```
Algorithm MergeSort(A)
begin
    mid = n/2; //divide step
    L = MergeSort(A[1..mid]);
    R = MergeSort(A[mid+1..n]);
    Merge(L, R); //combine step
end-Algorithm
```

- Algorithm mostly self-explanatory.

Divide and Conquer



- Example via merge sort
- Divide is split into two parts
- Recursively solve each subproblem

Runtime of Merge Sort

- Write the recurrence relation for merge sort and solve it.

Runtime of Merge Sort

- Write the recurrence relation for merge sort as $T(n) = 2T(n/2) + O(n)$.
 - This can be explained by the $O(n)$ time for merge and
 - The two subproblems obtained during the divide step each take $T(n/2)$ time.
 - Now use the general format for divide and conquer based algorithms.
- Solving this recurrence relation is done using say the substitution method giving us $T(n) = O(n \log n)$.
 - Look at previous examples.

Parallel Merge Sort

- An algorithm is a sequence of tasks T_1, T_2, \dots
- These tasks may have inter-dependencies,
 - Such as task T_i should be completed before task T_j for some i, j .
- However, it is often the case that there are several algorithms where many tasks are independent of each other.
 - In some cases, the algorithm or the computation has to be expressed in that independent-task fashion.
 - Example is parallel prefix.

Parallel Merge Sort

- In such a setting, one can imagine that tasks that are independent of each other can be done simultaneously, or in parallel.
- Let us think of arriving at a parallel merge sort algorithm.

Parallel Merge Sort

- What are the independent tasks in merge sort?
 - Sorting the two parts of the array.
 - This further breaks down to sorting four parts of the array, etc.
 - Eventually, every element of the array is a sorted sub-array.
 - So the main work is in merge itself.

Parallel Merge

- So, we just have to figure out a way to merge in parallel.
- Recall the merge algorithm as we developed it earlier.
 - Too many dependent tasks.
 - Not feasible in a parallel model.

Parallel Merge

- Need to rethink on a parallel merge algorithm
- Start from the beginning.
 - We have two sorted arrays L and R .
 - Need to merge them into a single sorted array A .
- Define the rank of an element x in a sorted array A as the number of elements of A that are smaller than x .
- To merge L and R , need to know the rank of every element from L and R in the merged array $L \cup R$.

Parallel Merge

- Importantly, for any x in L or R ,
$$\text{Rank}(x, L \cup R) = \text{Rank}(x, L) + \text{Rank}(x, R).$$
- So, merging is equivalent to finding the two ranks on the right hand side.

Parallel Merge

- Now, consider an element x in L at index k .
- How many elements of L are smaller than x ?
 - $k-1$.
- How many elements of R are smaller than x ?
 - No easy answer, but
 - can do binary search for x in R and get the answer.
 - Say k' elements in R are larger than x .

Parallel Merge

- How many elements in LUR are smaller than x ?
 - Precisely $k + k' - 1$.
- So, in the merged output, what index should x be placed in?
 - precisely at $k+k'$.
- Can this be done for every x in L ?
 - Yes, it is an independent operation.
- Can this be done for every x in R also?
 - Yes, replace the roles of L and R .
- All these operations are independent.

Example

$L = [8 \ 10 \ 12 \ 27]$

$R = [15 \ 17 \ 24 \ 32]$

Element	8	10	12	27	15	17	24	32
Rank in L	0	1	2	3	3	3	3	4
Rank in R	0	0	0	3	0	1	2	3
Rank in L U R	0	1	2	6	3	4	5	7

$LUR = [8 \ 10 \ 12 \ 15 \ 17 \ 24 \ 27 \ 32]$

Parallel Merge

- The above algorithm can be improved slightly.
- Need more techniques for that.
- So, it is a story left for another day.

Towards Parallel Sorting

- Use the parallel merge algorithm to sort.

```
Algorithm ParallelMergeSort(A)
Begin
    mid = n/2; //divide step
    L = MergeSort(A[1..mid]);
    R = MergeSort(A[mid+1..n]);
    Merge(L, R); //combine step
end-Algorithm
```